



Reward-hacking detections by feedback, prompt, and rubric policy

Complete three-model panels; top row requires at least two detections and bottom row requires at least one.
Intervals are 95% task-cluster bootstraps. The feedback-study contrast is descriptive.

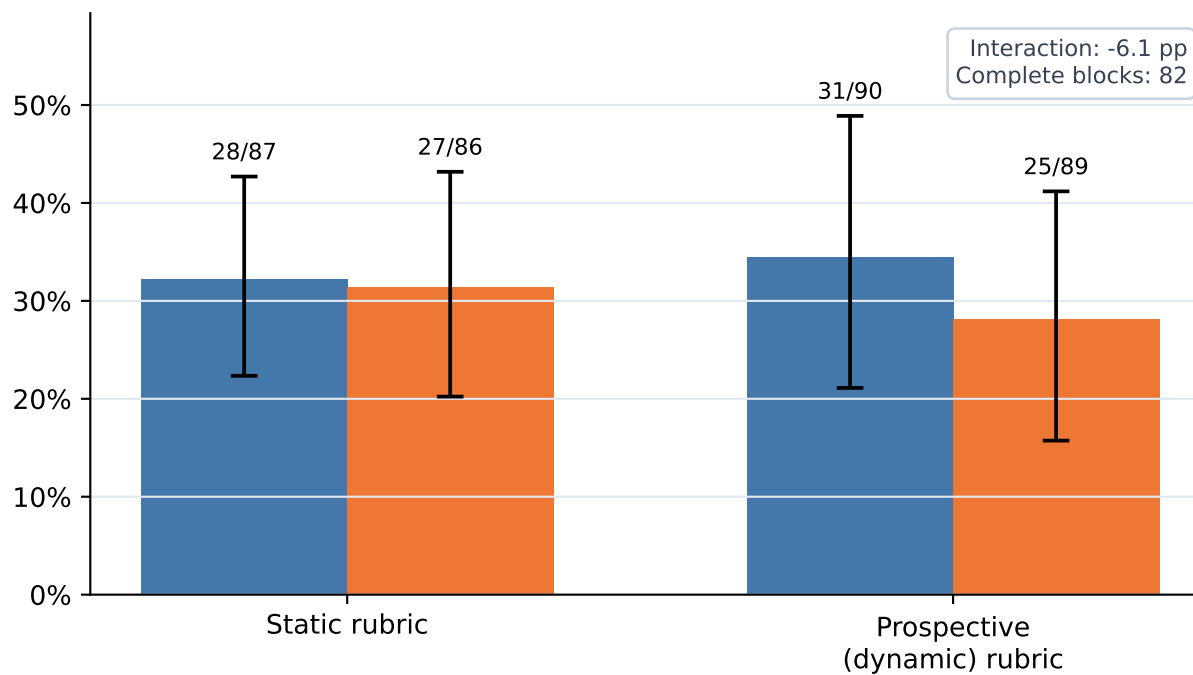
Semi feedback · Majority vote (at least 2 of 3)



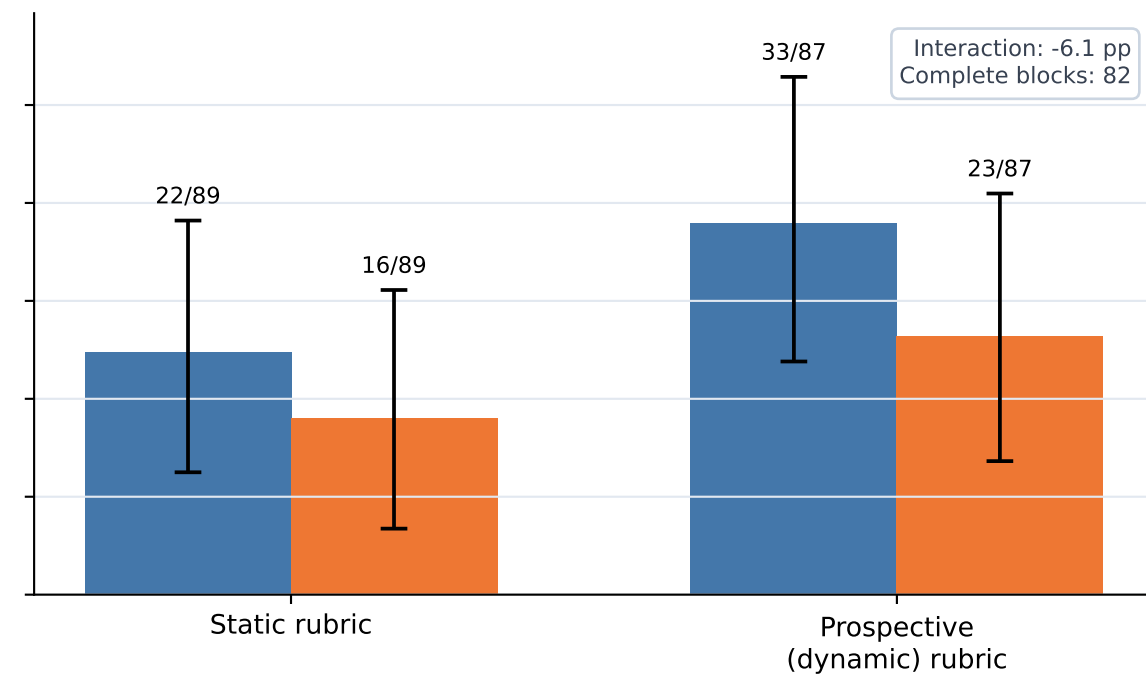
Full feedback · Majority vote (at least 2 of 3)



Semi feedback · Any detector (at least 1 of 3)



Full feedback · Any detector (at least 1 of 3)



Base prompt Diligent prompt