


Reward-hacking detections by feedback, prompt, and rubric policy

Three-model majority vote; 95% task-cluster bootstrap intervals. Feedback policies were run in separate studies, so their contrast is descriptive.


Semi feedback

Interaction: +6.1 pp
Complete blocks: 82



Full feedback

Interaction: -1.2 pp
Complete blocks: 82



Base prompt Diligent prompt